

```

import yfinance as yf
import pandas as pd

# Download data
apple_data = yf.download('AAPL', start='2021-01-01', end='2023-12-31',
interval='1d')
sp500_data = yf.download('^GSPC', start='2021-01-01', end='2023-12-
31', interval='1d')
sp500_it_data = yf.download('XLK', start='2021-01-01', end='2023-12-
31', interval='1d')

# Process the data
apple_data['Apple_Returns'] = apple_data['Adj Close'].pct_change()
sp500_data['SP500_Returns'] = sp500_data['Adj Close'].pct_change()
sp500_it_data['SP500_IT_Returns'] = sp500_it_data['Adj
Close'].pct_change()

# Combine the data into one dataframe
combined_data = pd.DataFrame({
    'Apple_Returns': apple_data['Apple_Returns'],
    'SP500_Returns': sp500_data['SP500_Returns'],
    'SP500_IT_Returns': sp500_it_data['SP500_IT_Returns']
})

# Drop missing values
combined_data.dropna(inplace=True)

# Save to a CSV
combined_data.to_csv('Cleaned_Combined_Data_2021_2023.csv')
print("\nCleaned and combined data saved as
'Cleaned_Combined_Data_2021_2023.csv'")

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

Cleaned and combined data saved as
'Cleaned_Combined_Data_2021_2023.csv'

combined_data = pd.read_csv('Cleaned_Combined_Data_2021_2023.csv',
index_col=0)
print(combined_data.head())

```

	Apple_Returns	SP500_Returns	SP500_IT_Returns
Date			
2021-01-05	0.012364	0.007083	0.006489
2021-01-06	-0.033661	0.005710	-0.017245
2021-01-07	0.034123	0.014847	0.026954

2021-01-08	0.008631	0.005492	0.006466
2021-01-11	-0.023249	-0.006555	-0.008259

```
# Calculate descriptive statistics
```

```
print("Descriptive Statistics:")
```

```
print(combined_data.describe())
```

```
Descriptive Statistics:
```

	Apple_Returns	SP500_Returns	SP500_IT_Returns
count	752.000000	752.000000	752.000000
mean	0.000705	0.000399	0.000697
std	0.017511	0.011083	0.015435
min	-0.058680	-0.043237	-0.053084
25%	-0.008891	-0.005863	-0.008470
50%	0.000776	0.000393	0.000582
75%	0.011262	0.006944	0.010097
max	0.088975	0.055434	0.082176

```
import matplotlib.pyplot as plt
```

```
# Use an available style
```

```
plt.style.use('ggplot')
```

```
# Plot time-series data
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(combined_data.index, combined_data['Apple_Returns'],  
label='Apple Returns', color='blue')
```

```
plt.plot(combined_data.index, combined_data['SP500_Returns'],  
label='S&P 500 Returns', color='orange')
```

```
plt.plot(combined_data.index, combined_data['SP500_IT_Returns'],  
label='S&P 500 IT Returns', color='green')
```

```
# Add titles and labels
```

```
plt.title("Daily Returns of Apple, S&P 500, and S&P 500 IT (2021-  
2023)", fontsize=14)
```

```
plt.xlabel("Date", fontsize=12)
```


```
plt.ylabel("Returns", fontsize=12)
```

```
# Add legend and grid
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```



```
import statsmodels.api as sm

# ===== Simple OLS Regression =====
# Independent variable: SP500 Returns
X_simple = combined_data['SP500_Returns']
# Dependent variable: Apple Returns
y = combined_data['Apple_Returns']

# Add a constant (intercept) to the model
X_simple = sm.add_constant(X_simple)

# Fit the regression model
simple_model = sm.OLS(y, X_simple).fit()

# Display the summary
print("\n===== Simple OLS Regression =====")
print(simple_model.summary())
```

```
===== Simple OLS Regression =====
                                OLS Regression Results

=====
Dep. Variable:                Apple_Returns    R-squared:
0.650
```

```

Model:                                OLS    Adj. R-squared:
0.649
Method:                               Least Squares    F-statistic:
1390.
Date:                                Mon, 16 Dec 2024    Prob (F-statistic):
6.24e-173
Time:                                13:59:08    Log-Likelihood:
2369.5
No. Observations:                     752    AIC:
-4735.
Df Residuals:                         750    BIC:
-4726.
Df Model:                             1

```

Covariance Type: nonrobust

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const              0.0002      0.000      0.519      0.604     -0.001
0.001
SP500_Returns      1.2733      0.034     37.284      0.000      1.206
1.340
=====
=====

```

```

=====
=====
Omnibus:                     62.545    Durbin-Watson:
1.845
Prob(Omnibus):               0.000    Jarque-Bera (JB):
329.299
Skew:                        0.010    Prob(JB):
3.12e-72
Kurtosis:                    6.242    Cond. No.
90.3
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

# Compute the correlation matrix
correlation_matrix = combined_data.corr()

```

```

# Display the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

```

Correlation Matrix:

	Apple_Returns	SP500_Returns	SP500_IT_Returns
Apple_Returns	1.000000	0.805943	0.876018
SP500_Returns	0.805943	1.000000	0.933453
SP500_IT_Returns	0.876018	0.933453	1.000000

```
import pandas as pd
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import
variance_inflation_factor

# Assuming combined_data is your processed dataset with returns
# Create the predictor matrix
X = combined_data[['SP500_Returns', 'SP500_IT_Returns']]

# Add a constant for the intercept
X = sm.add_constant(X)

# Calculate VIF for each predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = X.columns
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]

# Display VIF values
print("Variance Inflation Factor (VIF) for Predictors:")
print(vif_data)

Variance Inflation Factor (VIF) for Predictors:
      Variable      VIF
0          const  1.002338
1  SP500_Returns  7.772100
2  SP500_IT_Returns  7.772100

print("Columns in combined_data:", combined_data.columns)

Columns in combined_data: Index(['Apple_Returns', 'SP500_Returns',
'SP500_IT_Returns'], dtype='object')

# ===== Multiple OLS Regression =====
# Independent variables: SP500 Returns and SP500 IT Returns
X_multiple = combined_data[['SP500_Returns', 'SP500_IT_Returns']]
# Dependent variable: Apple Returns
y = combined_data['Apple_Returns']

# Add a constant (intercept) to the model
X_multiple = sm.add_constant(X_multiple)

# Fit the regression model
multiple_model = sm.OLS(y, X_multiple).fit()
```

```
# Display the summary
print("\n===== Multiple OLS Regression =====")
print(multiple_model.summary())
```

```
===== Multiple OLS Regression =====
                        OLS Regression Results
```

```
=====
```

```
Dep. Variable:          Apple>Returns    R-squared:
0.768
Model:                  OLS    Adj. R-squared:
0.768
Method:                 Least Squares    F-statistic:
1243.
Date:                  Mon, 16 Dec 2024    Prob (F-statistic):
1.08e-238
Time:                  13:59:21    Log-Likelihood:
2525.4
No. Observations:      752    AIC:
-5045.
Df Residuals:          749    BIC:
-5031.
Df Model:              2
```

```
Covariance Type:      nonrobust
```

```
=====
```

```
=====
```

	coef	std err	t	P> t	
[0.025	0.975]				

const	1.914e-06	0.000	0.006	0.995	-
0.001	0.001				
SP500>Returns	-0.1446	0.077	-1.868	0.062	-
0.297	0.007				
SP500_IT>Returns	1.0908	0.056	19.616	0.000	
0.982	1.200				

```
=====
```

```
=====
```

Omnibus:	43.059	Durbin-Watson:
1.916		
Prob(Omnibus):	0.000	Jarque-Bera (JB):
153.562		
Skew:	-0.054	Prob(JB):
4.51e-34		
Kurtosis:	5.211	Cond. No.
305.		

```
=====
=====

Notes:
```

```
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
```

```
# Residuals from the simple regression
```

```
residuals_simple = simple_model.resid
```

```
# Residuals from the multiple regression
```

```
residuals_multiple = multiple_model.resid
```

```
from scipy.stats import jarque_bera
```

```
jb_test_simple = jarque_bera(residuals_simple)
```

```
jb_test_multiple = jarque_bera(residuals_multiple)
```

```
print("Jarque-Bera Test for Simple Regression:", jb_test_simple)
```

```
print("Jarque-Bera Test for Multiple Regression:", jb_test_multiple)
```

```
Jarque-Bera Test for Simple Regression:
```

```
SignificanceResult(statistic=329.2988288682313,
```

```
pvalue=3.1165052487790823e-72)
```

```
Jarque-Bera Test for Multiple Regression:
```

```
SignificanceResult(statistic=153.56183080365957,
```

```
pvalue=4.513070579913583e-34)
```

```
from statsmodels.stats.stattools import durbin_watson
```

```
dw_stat_simple = durbin_watson(residuals_simple)
```

```
dw_stat_multiple = durbin_watson(residuals_multiple)
```

```
print("Durbin-Watson for Simple Regression:", dw_stat_simple)
```

```
print("Durbin-Watson for Multiple Regression:", dw_stat_multiple)
```

```
Durbin-Watson for Simple Regression: 1.8453927398290495
```

```
Durbin-Watson for Multiple Regression: 1.9158310974958823
```

```
from statsmodels.stats.api import het_white
```

```
# Robust covariance matrix
```

```
robust_cov = multiple_model.get_robustcov_results(cov_type="HAC",
maxlags=1)
```

```
print(robust_cov.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          Apple>Returns    R-squared:
0.768
```

```

Model:                                OLS    Adj. R-squared:
0.768
Method:                            Least Squares    F-statistic:
931.0
Date:                            Tue, 17 Dec 2024    Prob (F-statistic):
7.85e-204
Time:                            08:39:25    Log-Likelihood:
2525.4
No. Observations:                    752    AIC:
-5045.
Df Residuals:                        749    BIC:
-5031.
Df Model:                            2
Covariance Type:                    HAC

```

```

=====
=====
              coef    std err          t      P>|t|
[0.025    0.975]
-----
const                1.914e-06    0.000    0.006    0.995    -
0.001    0.001
SP500_Returns        -0.1446    0.111   -1.297    0.195    -
0.364    0.074
SP500_IT_Returns      1.0908    0.085   12.859    0.000
0.924    1.257
=====
=====
Omnibus:                43.059    Durbin-Watson:
1.916
Prob(Omnibus):          0.000    Jarque-Bera (JB):
153.562
Skew:                  -0.054    Prob(JB):
4.51e-34
Kurtosis:              5.211    Cond. No.
305.
=====
=====

```

Notes:

[1] Standard Errors are heteroscedasticity and autocorrelation robust (HAC) using 1 lags and without small sample correction

```
from arch import arch_model
```

```
# Fit GARCH(1,1) model
```

```
residuals_hac = multiple_model.resid # Residuals from the adjusted
regression
```



```
scaled_residuals_hac = residuals_hac * 100 # Scale residuals for  
better optimization
```

```
garch_model_hac = arch_model(scaled_residuals_hac, vol='Garch', p=1,  
q=1)
```

```
garch_results_hac = garch_model_hac.fit()
```

```
print(garch_results_hac.summary())
```

```
Iteration:      1,   Func. Count:      6,   Neg. LLF:
1715284.261092082
Iteration:      2,   Func. Count:     14,   Neg. LLF:
3095.590090631951
Iteration:      3,   Func. Count:     21,   Neg. LLF:
1183.411581073603
Iteration:      4,   Func. Count:     29,   Neg. LLF:
1370.5534682336702
Iteration:      5,   Func. Count:     35,   Neg. LLF:
1261.7497085208604
Iteration:      6,   Func. Count:     41,   Neg. LLF:
1224.8210608022048
Iteration:      7,   Func. Count:     47,   Neg. LLF:
1155.4260061277664
Iteration:      8,   Func. Count:     53,   Neg. LLF:
996.9167173594652
Iteration:      9,   Func. Count:     59,   Neg. LLF:
1029.6951125816108
Iteration:     10,   Func. Count:     65,   Neg. LLF:
984.9868031416943
Iteration:     11,   Func. Count:     71,   Neg. LLF:
959.7513084994887
Iteration:     12,   Func. Count:     77,   Neg. LLF: 938.672728439954
Iteration:     13,   Func. Count:     83,   Neg. LLF:
932.7711014256965
Iteration:     14,   Func. Count:     88,   Neg. LLF:
932.6920535666313
Iteration:     15,   Func. Count:     93,   Neg. LLF:
932.4747109494589
Iteration:     16,   Func. Count:     98,   Neg. LLF:
932.4225904846908
Iteration:     17,   Func. Count:    103,   Neg. LLF:
932.4201725766571
Iteration:     18,   Func. Count:    108,   Neg. LLF:
932.4183675894262
Iteration:     19,   Func. Count:    113,   Neg. LLF: 932.418176297936
Iteration:     20,   Func. Count:    118,   Neg. LLF:
932.4181455539419
Iteration:     21,   Func. Count:    122,   Neg. LLF:
932.4181455583173
Optimization terminated successfully      (Exit mode 0)
```

Current function value: 932.4181455539419
Iterations: 21
Function evaluations: 122
Gradient evaluations: 21
Constant Mean - GARCH Model Results

```
=====
=====
Dep. Variable:                None    R-squared:
0.000
Mean Model:                  Constant Mean    Adj. R-squared:
0.000
Vol Model:                   GARCH    Log-Likelihood:
-932.418
Distribution:                Normal    AIC:
1872.84
Method:                      Maximum Likelihood    BIC:
1891.33
                                No. Observations:
752
Date:                        Tue, Dec 17 2024    Df Residuals:
751
Time:                        08:57:31    Df Model:
1
```

Mean Model

```
=====
=====
              coef      std err          t      P>|t|      95.0%
Conf. Int.
-----
mu          2.0360e-03   3.027e-02   6.726e-02      0.946 [-5.729e-
02,6.137e-02]
```


Volatility Model

```
=====
=====
              coef      std err          t      P>|t|      95.0%
Conf. Int.
-----
omega        0.0257   1.184e-02     2.171   2.991e-02 [2.503e-
03,4.893e-02]
alpha[1]     1.7561e-03  1.903e-02   9.229e-02      0.926 [-3.554e-
02,3.905e-02]
beta[1]       0.9599   3.375e-02    28.443  5.998e-178 [ 0.894,
1.026]
```

Covariance estimator: robust

```
import matplotlib.pyplot as plt

# Plot the conditional volatility
conditional_volatility = garch_results_hac.conditional_volatility
plt.figure(figsize=(10, 6))
plt.plot(conditional_volatility, label="Conditional Volatility")
plt.title("GARCH(1,1) Conditional Volatility of Apple Returns")
plt.xlabel("Time")
plt.ylabel("Volatility")
plt.legend()
plt.show()
```



```
from statsmodels.graphics.tsaplots import plot_acf
import matplotlib.pyplot as plt

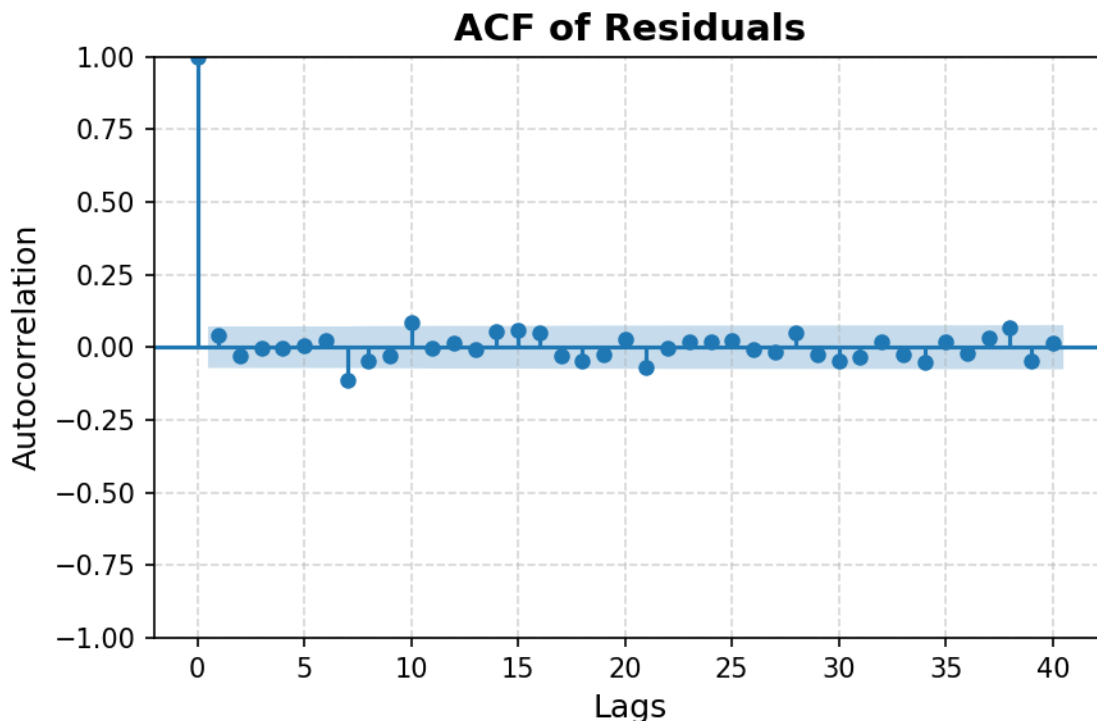
plt.style.use('default')

# Plot ACF
fig, ax = plt.subplots(figsize=(6, 4), dpi=150, facecolor='white')
plot_acf(residuals_hac, lags=40, alpha=0.05, title="ACF of Residuals",
color='tab:blue', ax=ax)
```

```

# Refine aesthetics
ax.set_facecolor("white") # Set axes background to white
plt.grid(True, linestyle='--', alpha=0.5)
plt.xlabel("Lags", fontsize=12)
plt.ylabel("Autocorrelation", fontsize=12)
plt.title("ACF of Residuals", fontsize=14, weight='bold')
plt.tight_layout()
plt.show()

```



```

from statsmodels.graphics.tsaplots import plot_pacf
import matplotlib.pyplot as plt


# Set plot style and figure size
plt.style.use('default') # Default clean style
fig, ax = plt.subplots(figsize=(6, 4), dpi=150, facecolor='white') #
White figure background

# Plot PACF with clean aesthetics
plot_pacf(residuals_hac, lags=40, alpha=0.05, title="PACF of
Residuals", color='tab:blue', ax=ax)

# Customize aesthetics
ax.set_facecolor("white") # Set axes background to white
plt.grid(True, linestyle='--', alpha=0.5) # Subtle gridlines
plt.xlabel("Lags", fontsize=12) # X-axis label

```

```
plt.ylabel("Partial Autocorrelation", fontsize=12) # Y-axis label
plt.title("PACF of Residuals", fontsize=14, weight='bold') # Improved title
plt.tight_layout()
plt.show()
```



```
from statsmodels.tsa.arima.model import ARIMA
```

```
# Fit ARIMA model to residuals
```

```
arima_model = ARIMA(residuals_hac, order=(1, 0, 1))
```

```
arima_results = arima_model.fit()
```

```
print(arima_results.summary())
```

```
C:\Users\shatakshi bansode\anaconda3\Lib\site-packages\statsmodels\
tsa\base\tsa_model.py:473: ValueWarning: A date index has been
provided, but it has no associated frequency information and so will
be ignored when e.g. forecasting.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\shatakshi bansode\anaconda3\Lib\site-packages\statsmodels\
tsa\base\tsa_model.py:473: ValueWarning: A date index has been
provided, but it has no associated frequency information and so will
be ignored when e.g. forecasting.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\shatakshi bansode\anaconda3\Lib\site-packages\statsmodels\
tsa\base\tsa_model.py:473: ValueWarning: A date index has been
```

provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
self._init_dates(dates, freq)

SARIMAX Results

=====					
=====					
Dep. Variable:	y	No. Observations:			
752					
Model:	ARIMA(1, 0, 1)	Log Likelihood			
2526.281					
Date:	Tue, 17 Dec 2024	AIC			
5044.562					
Time:	09:59:20	BIC			
5026.071					
Sample:	0	HQIC			
5037.438					
	- 752				
Covariance Type:	opg				
=====					
=====					
	coef	std err	z	P> z	[0.025
0.975]					

const	-9.934e-18	0.000	-3.14e-14	1.000	-0.001
0.001					
ar.L1	-0.4508	0.539	-0.836	0.403	-1.507
0.606					
ma.L1	0.4976	0.527	0.944	0.345	-0.536
1.531					
sigma2	7.069e-05	2.51e-06	28.129	0.000	6.58e-05
7.56e-05					
=====					
=====					
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):			
155.58					
Prob(Q):	0.91	Prob(JB):			
0.00					
Heteroskedasticity (H):	0.89	Skew:			
-0.05					
Prob(H) (two-sided):	0.36	Kurtosis:			
5.23					
=====					
=====					

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients  
(complex-step).
```

```
C:\Users\shatakshi bansode\anaconda3\Lib\site-packages\statsmodels\  
base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization  
failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to ")
```

```
# 1. OLS captures linear relationships.
```

```
# 2. ARIMA ensures that residuals are white noise.
```

```
# 3. GARCH captures volatility clustering, which is a key feature of  
financial returns.
```

```
# 4. Model selection was based on AIC, BIC, and volatility clustering.
```